

Evidence-based Integration Strategy

The following summarises the outsample evaluation (05 Jun 2025 – 05 Jun 2026; 1,770 games) of the V119 and V123 ensembles, individually and under the integration setups, and sets the staking policy for MLB Scion. All figures are flat 1-unit stakes (i.e. signal quality, before any Kelly / tiering is applied).

1. Strategies tested — performance summary

Strategy	Plays	% games	Win %	ROI (full)	ROI (2026)	Units	SR (per-bet)
V119 alone (line + stats)	357	20.2%	57.1%	10.3%	—	36.95	0.11
V123 alone (stats only)	526	29.7%	54.2%	9.5%	—	49.95	0.09
Setup 1 — Union (any signal)	572	32.3%	53.8%	6.7%	0.9%	38.49	0.07
Setup 2 — Consensus + V119-only	320	18.1%	58.8%	13.3%	8.9%	42.60	0.14
Setup 3 — Consensus + V123-only	444	25.1%	53.2%	7.4%	-0.2%	33.00	0.07
Setup 4 — Consensus ONLY	192	10.8%	60.4%	19.3%	12.0%	37.11	0.20

2. Most successful strategy — and why

Setup 4 (Consensus only) — play only when V119 and V123 independently back the same side — is the most successful strategy on every risk-adjusted measure, and the only one robust in the 2026 deployment regime.

It returns **19.3% ROI over the full outsample** (best-in-class Sharpe of 2.76 outsample / 0.20 per-bet) and **12.0% ROI in 2026 alone**, when the old program was losing on both sides. Consensus occurs in just 192 of 1,770 games (10.8%) — it is selective, not high-volume. The evidence is unambiguous that **selectivity is the edge**: playing every signal (Union) is the worst outcome (6.7% full, 0.9% in 2026), and each step away from pure consensus lowers ROI.

3. Why two ensembles beat one — information-content decomposition

Every setup is built from three disjoint buckets. Isolating them shows exactly where the edge lives (counts reconcile to the Union: $192 + 128 + 252 = 572$).

Information bucket	Plays	Win %	ROI (full)	ROI (2026)	Dominant side → its ROI
(i) V119 & V123 AGREE (consensus)	192	60.4%	+19.3%	+12.0%	Both strong; 2026 dogs +19.1%, faves +1.4%
(ii) V119 plays, V123 silent	128	56.3%	+4.3%	+4.5%	69% faves → +4.8% (2026 faves +7.3%)
(iii) V119 silent, V123 plays	252	47.6%	-1.6%	-9.3%	83% dogs → +0.6% full / -10.3% (2026)

Where the two ensembles **agree**, ROI roughly doubles (from ~10% alone to 19.3%) — because a market-aware model (V119) and a market-blind model (V123) converging is two independent lines of evidence, not one. That agreement is high-conviction information neither model can produce alone.

4. Staking policy — where to stake, what to ignore

Stake level	Which bets	Why (evidence)
FULL stake	Consensus plays — both ensembles agree (priority: consensus DOGS)	The 19.3% full / 12.0% (2026) engine. Consensus dogs are robust across both periods (+18.8% / +19.1%).
HALF stake	(a) Consensus FAVOURITES, in the current regime; (b) optionally, V119-only plays (V119 backs, V123 silent)	(a) Consensus faves were +20.3% full but fell to +1.4% in 2026 — regime risk, so de-risk until the fave edge recovers. (b) V119-only is mildly positive (+4.3% / +4.5%) but is the least-stable bucket (only 15 V119 models) and mostly faves.
DO NOT PLAY	V123-only plays (V123 backs, V119 silent). Also: never play the Union.	No durable edge: solo dogs +0.6% full but -10.3% in 2026; solo faves -12.7%; bucket -1.6% / -9.3%. The Union is the worst overall strategy (6.7% / 0.9%).

In short: stake the agreement, halve the favourite / single-model exposure, and never let V123 play alone.

5. Strengths and weaknesses of each ensemble

V119 (line + stats). Includes the opening line, the bet line, the difference between bet and opening line, and the average historical bookie probabilities for the Home and Visitor teams over each team's last X home / visitor games respectively.

Strengths: the line features carry genuine market information, giving V119 real standalone edge — concentrated on **favourites**, whose pricing is line-driven (its solo favourites earn +4.8% full / +7.3% in 2026). Highest win rate of any single model (57.1%). It is the primary directional signal.

Weaknesses: good V119 models are scarce (only 15, vs V123's 41) — partly because it competes with the market it ingests, so beating the line while using the line is a narrow target. Its 10.3% standalone ROI is flattered by the consensus games; stripped of those, its independent edge is modest (+4.3%).

V123 (stats only). The same stat features as V119 but **without** the opening and bet lines for the game being predicted (and therefore without the bet–open line difference). Their place is taken by other highly ranked statistics (e.g. BullpenERA).

Strengths: being market-blind, its agreement with V119 is genuinely **independent** confirmation — and that confirmation is what creates the consensus edge (ROI roughly doubles). Good models are easier to find (41). It is a powerful confirmer and the only ensemble that can diverge from the market.

Weaknesses: no durable standalone edge. Its solo plays do not profit — solo dogs +0.6% full but **-10.3% in 2026**; solo favourites -12.7%. Without the line it cannot price favourites, and its stat-derived “value” on dogs evaporates out-of-regime. **It must never play alone.**